


GOOGLE CLOUD PLATFORM

CLOUD SPANNER





WHAT IS CLOUD SPANNER?

Description


- Mission-critical, relational database service with transactional consistency, global scale and high availability.

Good for

- Mission-critical applications
- High transactions
- Scale + Consistency requirements


Common Workloads

- Adtech
- Financial services
- Global supply chain
- Retail




FEATURES

- Fully managed database service with a global scale
- Traditional relational semantics: schemas, ACID transactions, SQL
- Automate, synchronous replication within/across regions for availability, maintained by site reliability engineers
- Uses nodes for scalability - each node provides up to 2 TiB of storage
- No backup solution available
- Used at Google for Google Play, Adwords, etc



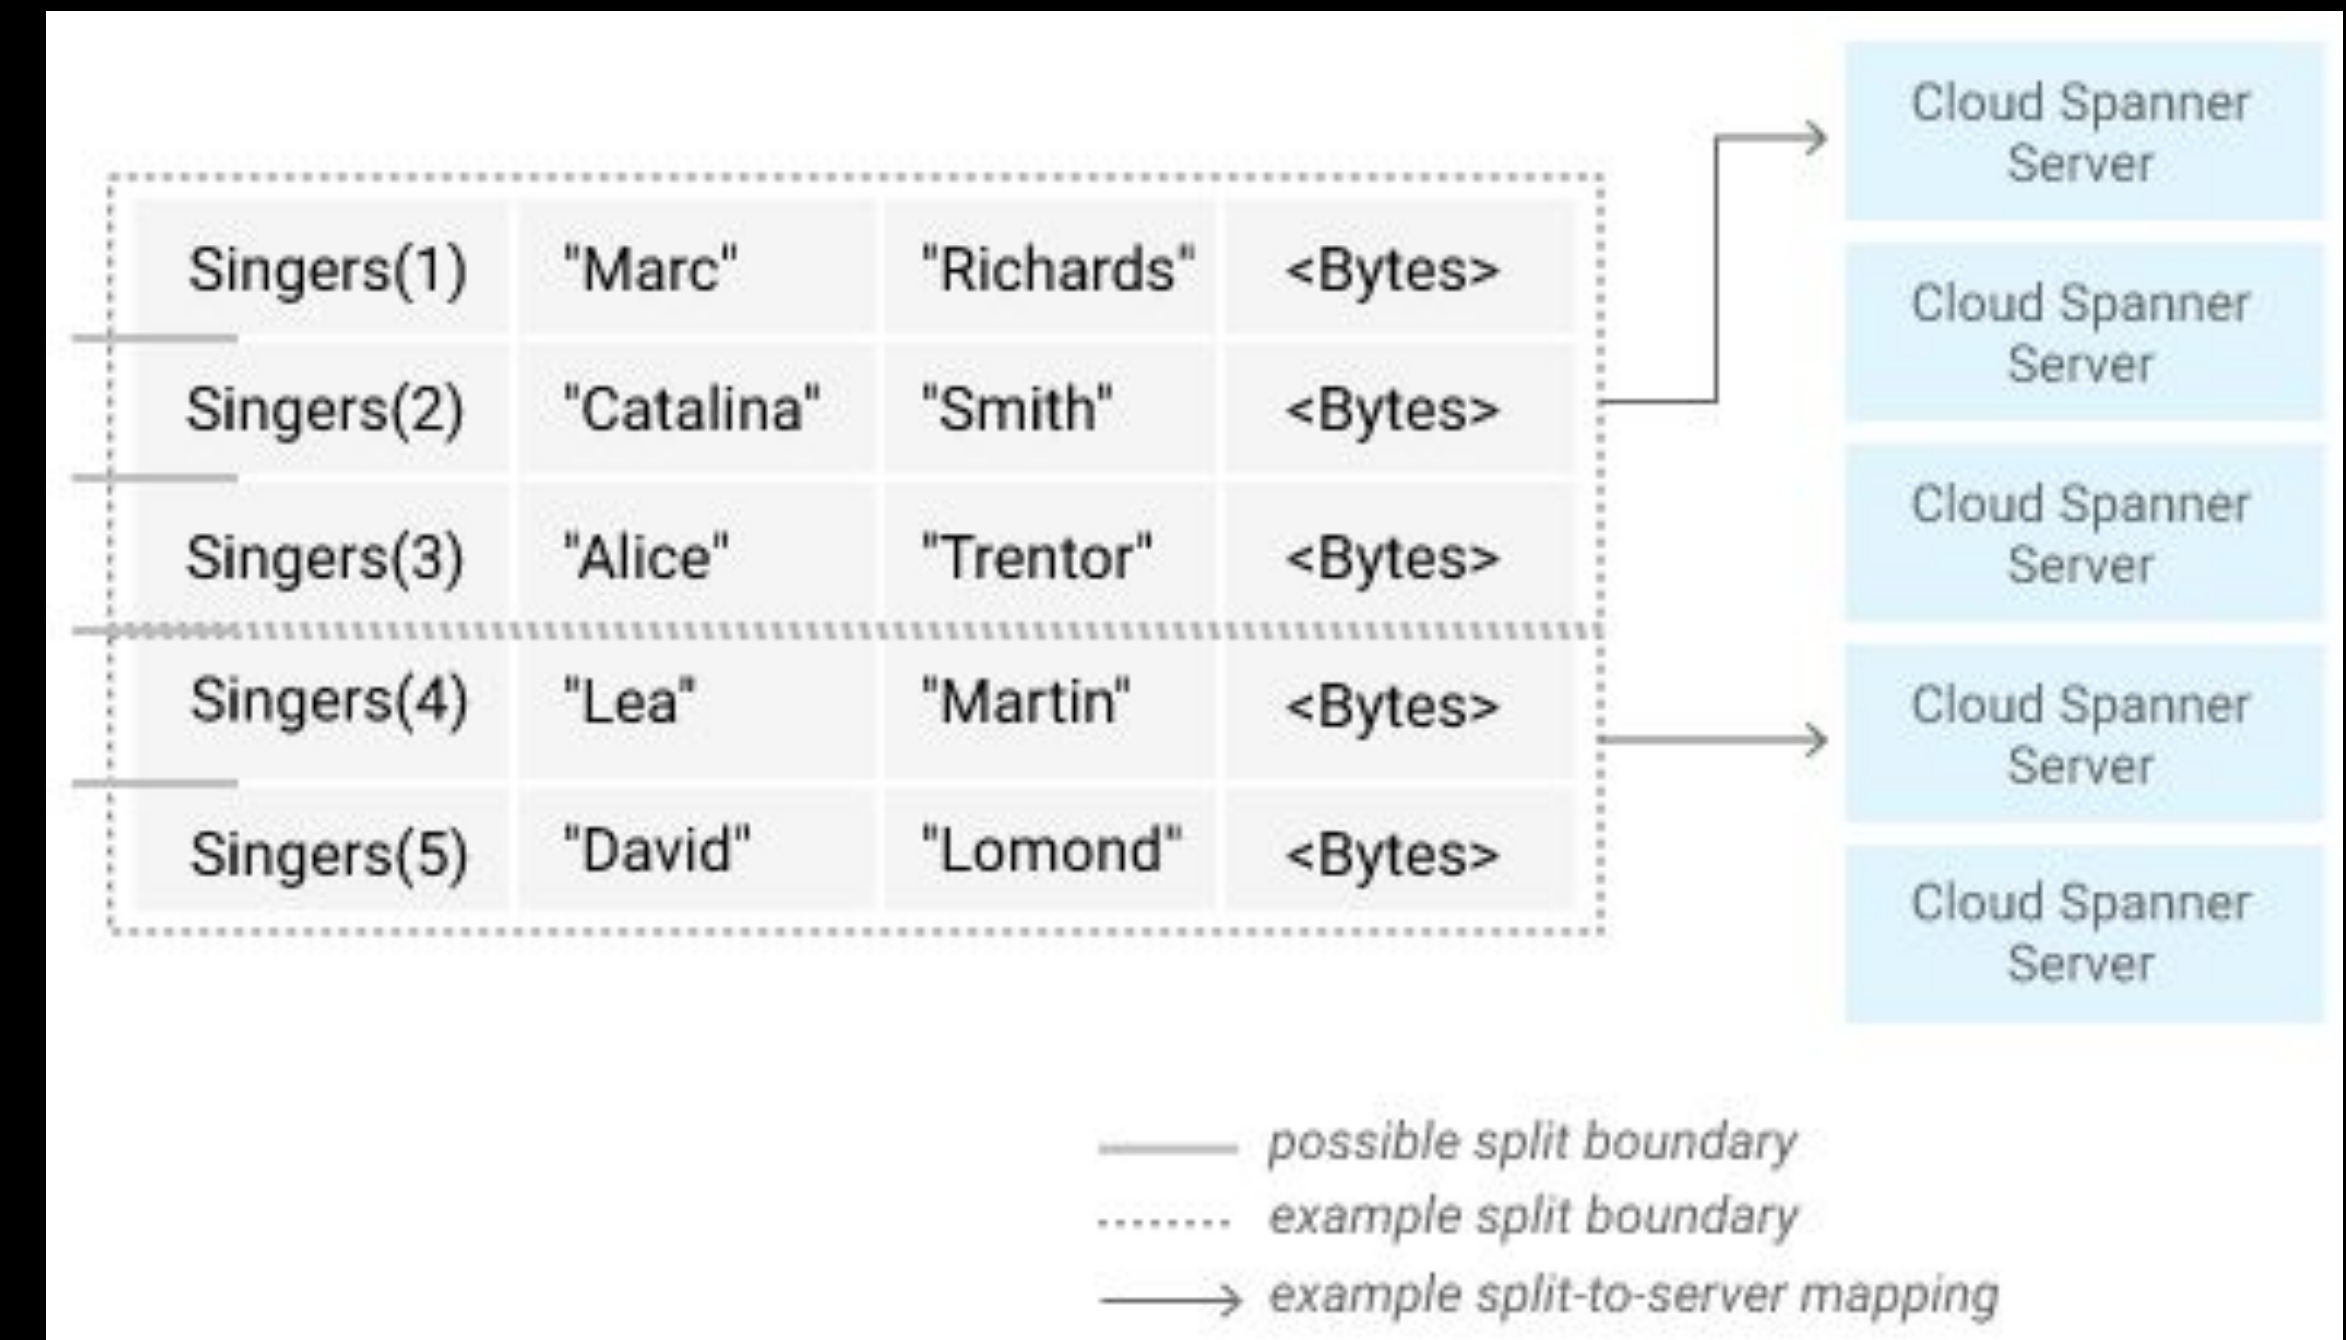
TRADITIONAL VS. SPANNER


	SPANNER	TRADITIONAL RELATIONAL	TRADITIONAL NON-RELATIONAL
Schema	Yes	Yes	No
SQL	Yes	Yes	No
Consistency	Strong	Strong	Eventual
Availability	High	Failover	High
Scalability	Horizontal	Vertical	Horizontal
Replication	Automatic	Configurable	Configurable



SCHEMA AND DATA MODEL

- Tables look like relational database tables in that they are structured with rows, columns, and values, and they contain primary keys.
- Data is strongly typed: you must define a schema for each database and that schema must specify the data types of each column of each table
- Cloud Spanner divides your data into chunks called "splits" which can move to different nodes/servers






INTERLEAVE DATA

- An **interleaved table** is a table that you declare to be a child of another table because you want the rows of the child table to be physically stored together with the associated parent row.
- E.g. Three tables – Singers, Albums and Songs
- Physically stored as one table, but gets split across nodes


Singers(1)	"Marc"	"Richards"	<Bytes>		
Albums(1, 1)				"Total Junk"	
Albums(1, 2)				"Go, Go, Go"	
Songs(1, 2, 1)					"42"
Songs(1, 2, 2)					"Nothing Is The Same"
Singers(2)	"Catalina"	"Smith"	<Bytes>		
Albums(2, 1)				"Green"	
Songs(2, 1, 1)					"Let's Get Back Together"
Songs(2, 1, 2)					"Starting Again"
Songs(2, 1, 3)					"I Knew You Were Magic"
Albums(2, 2)				"Forever Hold Your Peace"	
Albums(2, 3)				"Terrified"	
Songs(2, 3, 1)					"Fight Story"

— possible split boundary




REPLICATION

- Cloud Spanner creates replicas of each database split
- Three types: Read-write, Read-only, Witness
- Certain parts of the data in the replicas is owned by different nodes – custom algorithm automatically manages data
- Benefits includes Data availability, Geographic locality, Single database experience and Easier application development

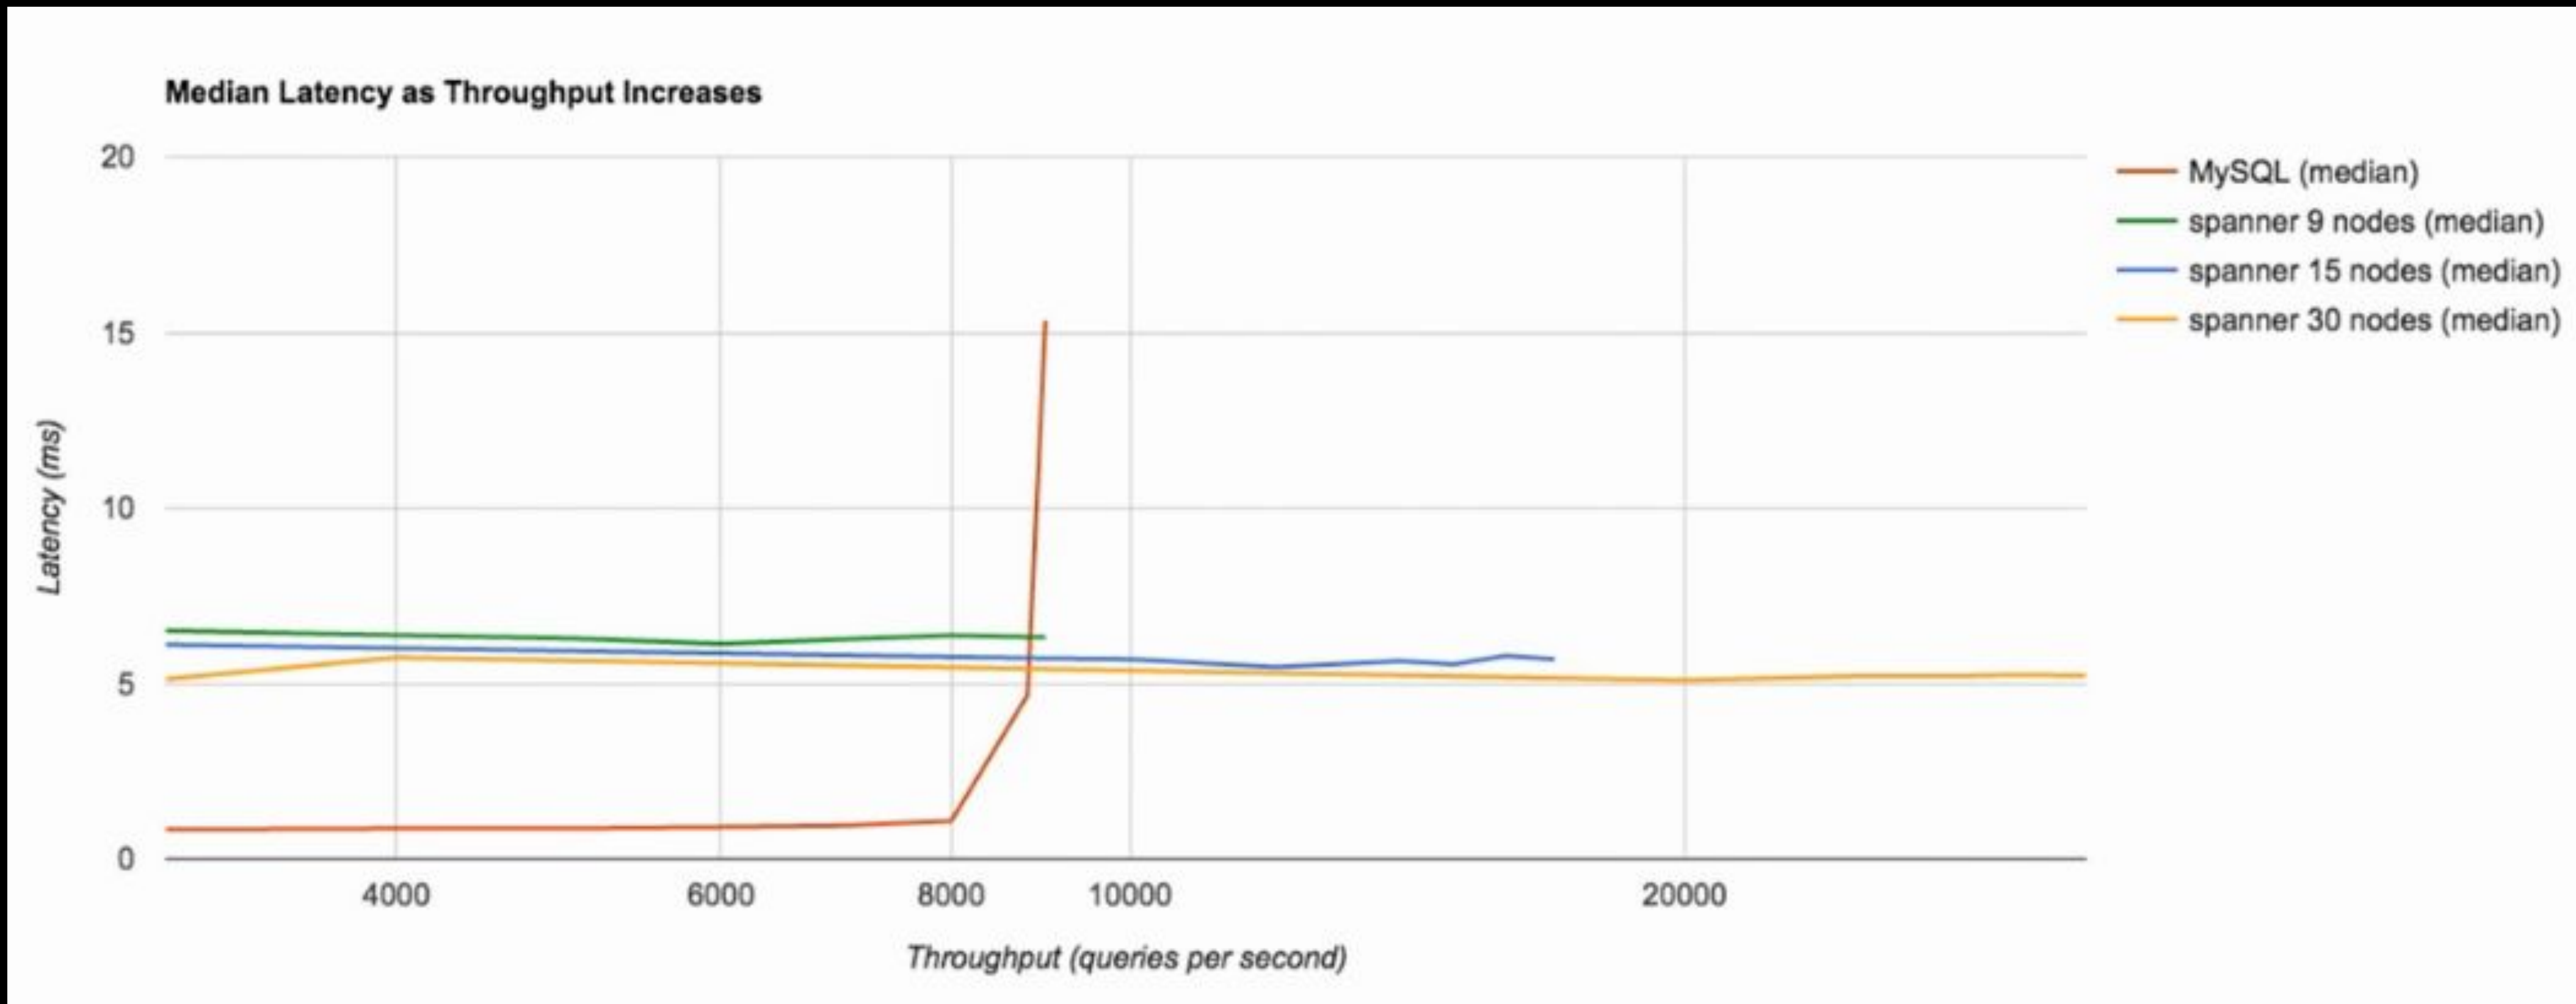


DEMO: CREATE CLOUD SPANNER


1. Select/create project
2. Navigate to the Cloud Spanner dashboard
3. Create Instance
4. Create a new database
5. Create two interleaved tables



HORIZONTAL SCALABILITY



- Increase the number of nodes and splits are automatically optimized across replicas over regions (Quizelet app) – 700 GB, 6 Billion rows



TRUE TIME & EXTERNAL CONSISTENCY

- TrueTime is a distributed clock that enables applications to generate monotonically increasing timestamps
- Cloud spanner uses this to timestamp transactions, allowing it to perform consistent reads across an entire database and across multiple Cloud regions without blocking writes.
- With **External Consistency**, the system behaves as if all transactions were executed sequentially, even though Cloud Spanner actually runs them across multiple servers
- External consistency is a stronger property than *strong consistency* as it doesn't block data during *strong reads*